

DATA 514 Section 9 Worksheet - Jonathan Jacobs

Question 1: Wide Column Database Implementation

Keyspace and Column Families

Keyspace: `ParkingEnforcement`

Column Families:

1. **Payroll**
2. **Regist**
3. **ParkingTickets**

Detailed Schema

1. **Payroll:**

- **Row Key:** `UserID`
- **Columns:**
 - `Name`
 - `Job`
 - `ParkingPermit`

2. **Regist:**

- **Row Key:** `UserID`
- **Columns:**
 - `Car_<index>` (for handling multiple cars, e.g., `Car_1`, `Car_2`)
 - `LicensePlate_<index>` (for handling multiple license plates, e.g., `LicensePlate_1`, `LicensePlate_2`)

3. **ParkingTickets:**

- **Row Key:** `LicensePlate`
- **Columns:**
 - `ParkingLot_<index>` (for handling multiple tickets, e.g., `ParkingLot_1`, `ParkingLot_2`)
 - `Date_<index>` (for handling multiple dates, e.g., `Date_1`, `Date_2`)
 - `Amount_<index>` (for handling multiple amounts, e.g., `Amount_1`, `Amount_2`)

Implementation Details

Payroll:

- The `Payroll` column family stores details about the employees. Each row is keyed by the `UserID`, ensuring unique entries for each user.
- Columns within each row include the `Name`, `Job`, and `ParkingPermit` of the employee. **Regist:**
- The `Regist` column family stores information about the cars owned by users. Each row is keyed by the `UserID`, ensuring unique entries for each user.

- Users can own multiple cars, so the schema supports multiple entries using indexed columns. For example, a user with two cars would have columns **Car_1**, **LicensePlate_1**, **Car_2**, and **LicensePlate_2**. **ParkingTickets:**
- The **ParkingTickets** column family tracks parking tickets. Each row is keyed by the **LicensePlate**, ensuring unique entries for each car.
- A single car can have multiple tickets, so the schema supports multiple entries using indexed columns. For example, a car with two tickets would have columns **ParkingLot_1**, **Date_1**, **Amount_1**, **ParkingLot_2**, **Date_2**, and **Amount_2**.

Use of Explicit Timestamps

- Explicit timestamps can be used to track the creation and modification times of each column entry, ensuring data consistency and enabling versioning if necessary.
- In a wide column database like Cassandra, each column can have an associated timestamp that indicates when the data was written, which is useful for handling conflicts and ensuring data integrity.

Question 2: Graph Database

Nodes and Relationships

Nodes:

1. **User**
2. **Car**
3. **ParkingTicket**

Relationships:

1. **OWNED_BY**
2. **HAS_TICKET**

Node Details

1. **User:**

- Properties:
 - **UserID**
 - **Name**
 - **Job**
 - **ParkingPermit**

2. **Car:**

- Properties:
 - **LicensePlate**
 - **Model**

3. **ParkingTicket:**

- Properties:
 - **TicketID**
 - **Date**
 - **Amount**
 - **ParkingLot**

Relationship Details

1. OWNED_BY:

- Connects: **Car** -> **User**
- Properties: None

2. HAS_TICKET:

- Connects: **Car** -> **ParkingTicket**
- Properties: None

Implementation Details

User Nodes:

- Represent individuals in the payroll system.
- Each **User** node includes properties for **UserID**, **Name**, **Job**, and **ParkingPermit**.
- Represent cars registered in the system.
- Each **Car** node includes properties for **LicensePlate** and **Model**.
- Represent individual parking tickets.
- Each **ParkingTicket** node includes properties for **TicketID**, **Date**, **Amount**, and **ParkingLot**.

Relationships:

- **OWNED_BY**: This relationship connects a **Car** node to a **User** node, indicating ownership. Each car is related to its owner via this relationship.
- **HAS_TICKET**: This relationship connects a **Car** node to a **ParkingTicket** node, indicating that the car has received a particular parking ticket. Each car can have multiple **HAS_TICKET** relationships, each linking to a different parking ticket.

Queries Supported

1. Listing the permitted parking lot and per-car tickets incurred by each user:

- Traverse from **User** to **Car** via **OWNED_BY** relationships, and from **Car** to **ParkingTicket** via **HAS_TICKET** relationships to list all tickets for each user's cars.

2. Counting how many tickets a license plate has ever had:

- Traverse from **Car** to **ParkingTicket** via **HAS_TICKET** relationships and count the number of connections.

3. Determining whether a plate is allowed to be in a specific lot:

- Check the **ParkingPermit** property of the **User** node connected to the **Car** node via **OWNED_BY** relationship to verify if the car is allowed in the specified parking lot.

Graph Database Schema in Mermaid Notation

```
erDiagram
    User {
        int UserID
        string Name
        string Job
        string ParkingPermit
    }
    Car {
```

```
        string LicensePlate
        string Model
    }
    ParkingTicket {
        int TicketID
        date Date
        string ParkingLot
        float Amount
    }
    User ||--o{ Car : owns
    Car ||--o{ ParkingTicket : has_ticket
```